

ECE 158B Data Networks II

Lecture 09. P2P; DHT; CDN

Prof. Xinyu Zhang

<http://xyzhang.ucsd.edu>

Department of Electrical and Computer Engineering

University of California San Diego



Today's agenda

- P2P architecture and applications
- DHT: Distributed Hash Table
- CDN

P2P architecture and applications


P2P architecture

➤ Characteristics

- no always-on server
- arbitrary end systems directly communicate
- peers request service from other peers, provide service in return to other peers.
Self scalability – new peers bring new service capacity, and new service demands
- peers are intermittently connected and change IP addresses – complicating management

➤ Examples:

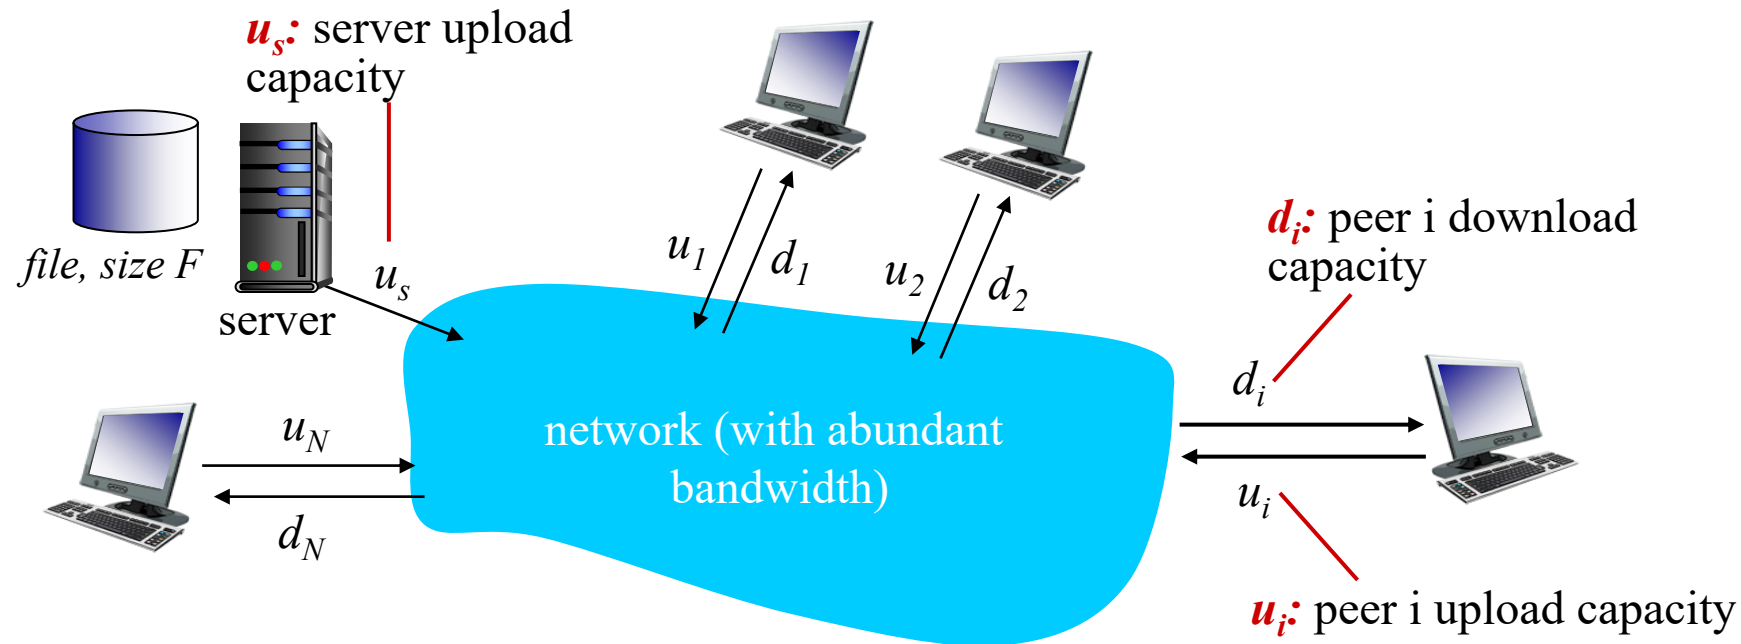
- File distribution (BitTorrent), Streaming (PPLive), VoIP (Skype)



Client/server vs. P2P architecture

Question: how much time to distribute file (size F) from one server to N peers?


- peer upload/download capacity is limited resource



Client/server vs. P2P architecture

File distribution time using client-server

- **server transmission:** must sequentially send (upload) N file copies:
 - ✓ time to send one copy: F/u_s
 - ✓ time to send N copies: NF/u_s
- **client:** each client must download file copy
 - ✓ d_{\min} = min client download rate
 - ✓ max client download time: F/d_{\min}



time to distribute F
to N clients using
client-server approach


$$D_{c-s} \geq \max\{NF/u_s, F/d_{\min}\}$$

increases linearly in N

Client/server vs. P2P architecture

File distribution time using P2P

- **server transmission:** must upload at least one copy
 - ✓ time to send one copy: F/u_s
- **client:** each client must download file copy
 - ✓ max client download time: F/d_{\min}
- **clients:** as aggregate must download NF bits
 - max upload rate (limiting max download rate) is $u_s + \sum u_i$



time to distribute F
to N clients using
P2P approach


$$D_{P2P} \geq \max\{F/u_s, F/d_{\min}, NF/(u_s + \sum u_i)\}$$

increases linearly in N ...

... but so does this, as each peer brings service capacity

Client/server vs. P2P architecture


client upload rate = u , $F/u = 1$ hour, $u_s = 10u$, $d_{min} \geq u_s$



P2P file distribution protocol: BitTorrent

- file divided into 256Kb chunks
- peers in torrent send/receive file chunks

tracker: tracks peers
participating in torrent




Alice arrives ...
... obtains list
of peers from tracker
... and begins exchanging
file chunks with peers in torrent

torrent: group of peers
exchanging chunks of a file

BitTorrent: initialization, and dynamics

- New peer joining torrent:
 - has no chunks, but will accumulate them over time from other peers
 - registers with tracker to get list of peers, connects to subset of peers (“neighbors”)
- while downloading, peer uploads chunks to other peers
- peer may change peers with whom it exchanges chunks
- once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent
 - Peers may come and go (referred to as “churn”)



BitTorrent: requesting and sending file chunks

Chunk selection: rarest first


- at any given time, different peers have different subsets of file chunks
- periodically, Alice asks each peer for list of chunks that they have
- Alice requests missing chunks from peers, **rarest first**

Peer selection: tit-for-tat with optimistic unchoke

- Alice sends chunks to those 4 peers currently sending her chunks *at highest rate*
 - ✓ other peers are **choked** by Alice (do not receive chunks from her)
 - ✓ re-evaluate top 4 every 10 seconds
- Exploring new peers: randomly select another peer every 30 seconds, starts sending chunks
 - ✓ “optimistically unchoke” this peer
 - ✓ newly chosen peer may join top 4

BitTorrent peer selection example

- (1) Alice “optimistically unchokes” Bob
- (2) Alice becomes one of Bob’s top 4 providers; Bob reciprocates
- (3) Bob may become one of Alice’s top 4 providers as well



Can tit-for-tat completely remove free-riders?

DHT: another P2P application

- DHT: a distributed P2P database
- Database has (key, value) pairs; examples:
 - key: SSN; value: human name
 - key: movie title; value: IP address
- DHT: Distribute the (key, value) pairs over the (millions of peers)
 - a peer queries DHT with key
 - DHT returns values that match the key
 - peers can also insert (key, value) pairs

DHT: another P2P application

- Motivation of building a distributed database
 - Realizing content addressable network
 - Scalability: No peer is powerful enough to keep a million-entry database
- Use cases: real-world P2P file-sharing or P2P multimedia systems
 - e.g., BitTorrent can use a DHT to create a distributed tracker.
 - ✓ Key: the torrent identifier
Value: IP addresses of all the peers currently participating in the torrent

How to assign (key, value) entries to peers?

- The central issue of DHT!
- Basic idea:
 - Convert each key to an integer
 - Assign integer “ID” to each peer
 - Put (key,value) pair in the peer whose integer ID is closest to the key

How to assign (key, value) to peers?

- assign integer identifier to each peer in range $[0, 2^n - 1]$ for some n
 - each identifier represented by n bits.
- require each key to be an integer in same range
 - to get integer key, hash original key
 - ✓ e.g., $\text{key} = \text{hash}(\text{"Led Zeppelin IV"})$
 - This is why it is referred to as a distributed "hash" table (DHT)


A **hash function** is a many-to-one function for which two different inputs can have the same output (same integer), but the likelihood of having the same output is extremely small.

Think of it as a **function that compresses information into a few digits.**

How to assign (key, value) to peers?

- rule: assign key to the peer that has the **closest** ID.
 - Closest: equals the key, or immediate successor of the key
 - e.g., $n=4$; peers: 1,3,4,5,8,10,12,14
 - ✓ key = 13, then successor peer = 14
 - ✓ key = 15, then successor peer = 1 (modulo operation)


(key, value) assignment in a circular DHT




- Each peer *only* aware of immediate successor and predecessor.
- An application layer “overlay network”

(key, value) assignment in a circular DHT

$O(N)$ messages
on average to resolve
query, when there
are N peers



(key, value) assignment in a circular DHT, with shortcuts



- each peer keeps track of IP addresses of predecessor, successor, short cuts.
- reduced from 6 to 2 messages.
- Proven result: possible to design shortcuts so $O(\log N)$ neighbors, $O(\log N)$ messages in query


Handling peer churns in circular DHT

➤ How to handle peer churns?

- Each peer knows address of its two successors
- Each peer periodically pings its two successors to check aliveness
- If immediate successor leaves, choose next successor as new immediate successor

➤ Example: peer 5 abruptly leaves

- Peer 4 detects peer 5 departure; makes 8 its immediate successor; asks 8 who its immediate successor is; makes 8's immediate successor its second successor.
- What if peer 13 wants to join, and it only knows peer 1?



Content distribution networks (CDN)


Why CDN?

- Challenge: how to stream content (selected from millions of videos) to hundreds of thousands of simultaneous users?
- Potential solution: single, large “mega-server”?
 - single point of failure
 - point of network congestion
 - long (and possibly congested) path to distant clients

Centralized solution doesn't scale well!

Why CDN?

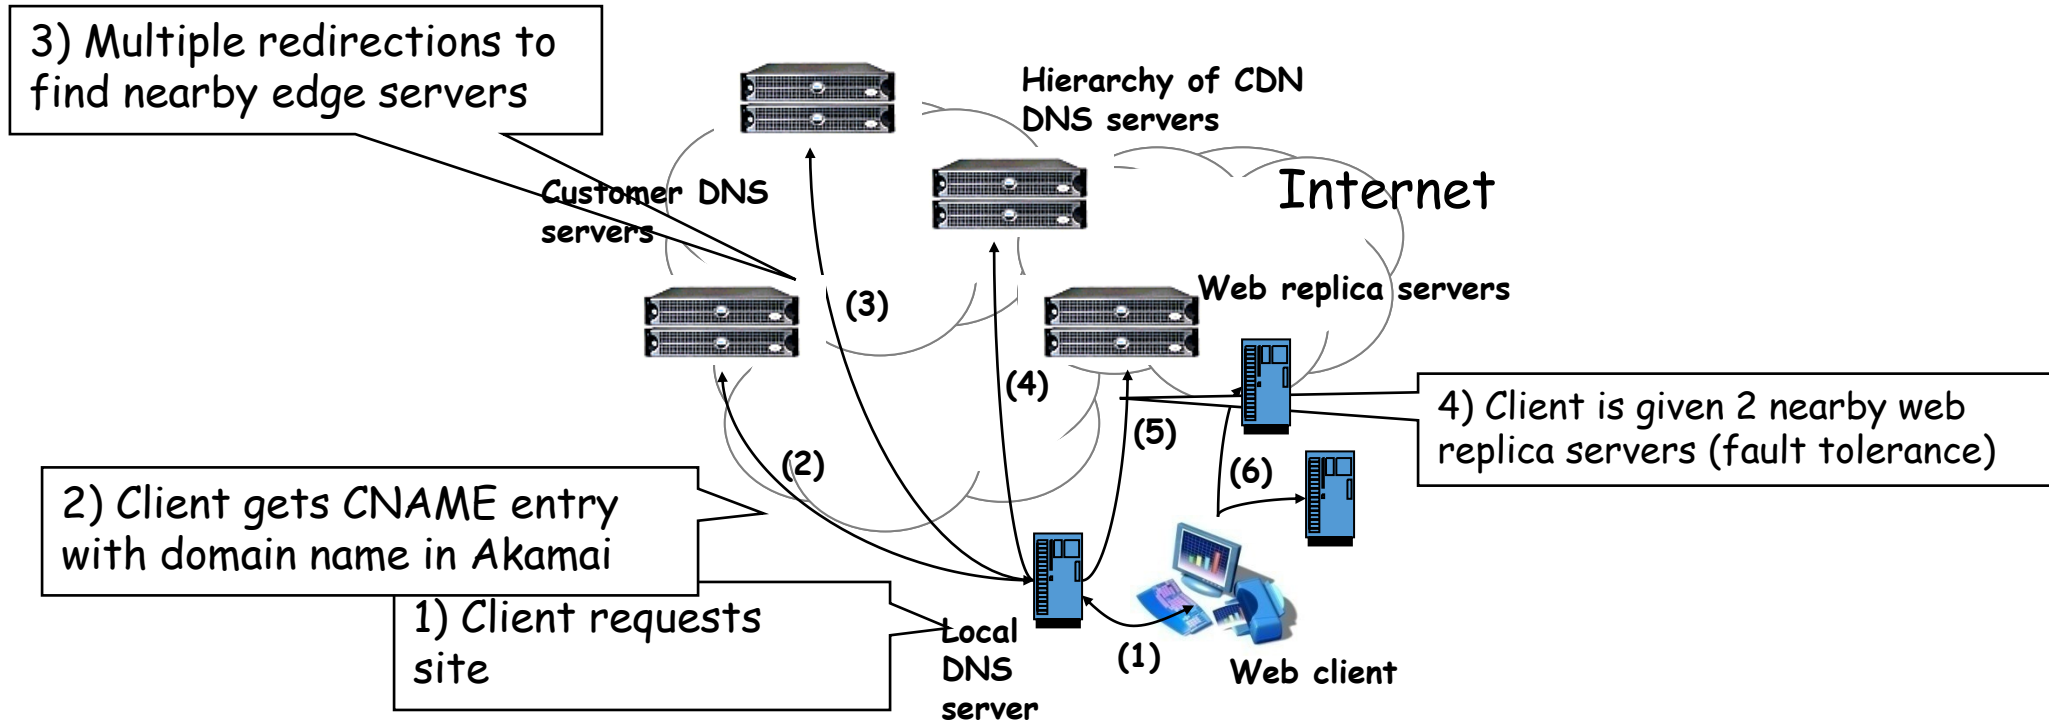
- Recap: web cache
- Can web cache satisfy needs of all content publishers? e.g., YouTube
 - Freshness of content
 - Performance scalability (multiple servers)
 - Flexibility
 - Direct control
- CDN is thus designed, e.g.,
 - Akamai, YouTube CDN




Basics of Akamai CDN

- Content publisher (e.g., CNN, NYTimes)
 - provides base HTML documents
 - runs origin server(s)
- Akamai
 - Runs edge servers for hosting content; deep deployment into 1000+ networks
 - Direct control
- customized DNS redirection servers to select edge servers based on
 - closeness to client browser
 - server load

Akamai load direction flow



Akamai today



Source: <https://networkingchannel.eu/living-on-the-edge-for-a-quarter-century-an-akamai-retrospective-downloads/>

Review questions

- ✓ P2P: quantitative understanding of the advantage of P2P file distribution in comparison to client/server; Important operations of BitTorrent, e.g., chunk request, chunk selection, peer selection, optimistic unchoking, tit-for-tat
- ✓ DHT: motivation and use cases; assignment of (key, value) to a circular DHT; handling peer churns
- ✓ Web caching vs. CDN (Why is CDN needed if we already have Web cache?)

References

- ✓ Chapter 2.6 (Peer-to-Peer Applications) of the book “Computer networking : a top-down approach”
- ✓ Chapter 7.2.4 (Content Distribution Networks) of the book “Computer networking : a top-down approach”